References (delete or hide for presentation)

- Rubric with tasks - link
- Team project plan - link
- Roadmap - link

MIDSearch

A RAG based multiplatform search tool for the MIDS program

Devin Suy, Nadia Tantsyura, Randy Louie, Robert Greer, Thomas Lai

Team

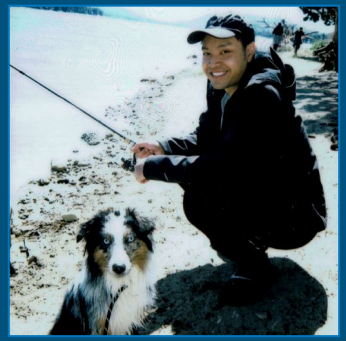

Devin Suy

Nadia Tantsyura

Thomas Lai

Randy Louie

Robert Greer

Problem & Motivation

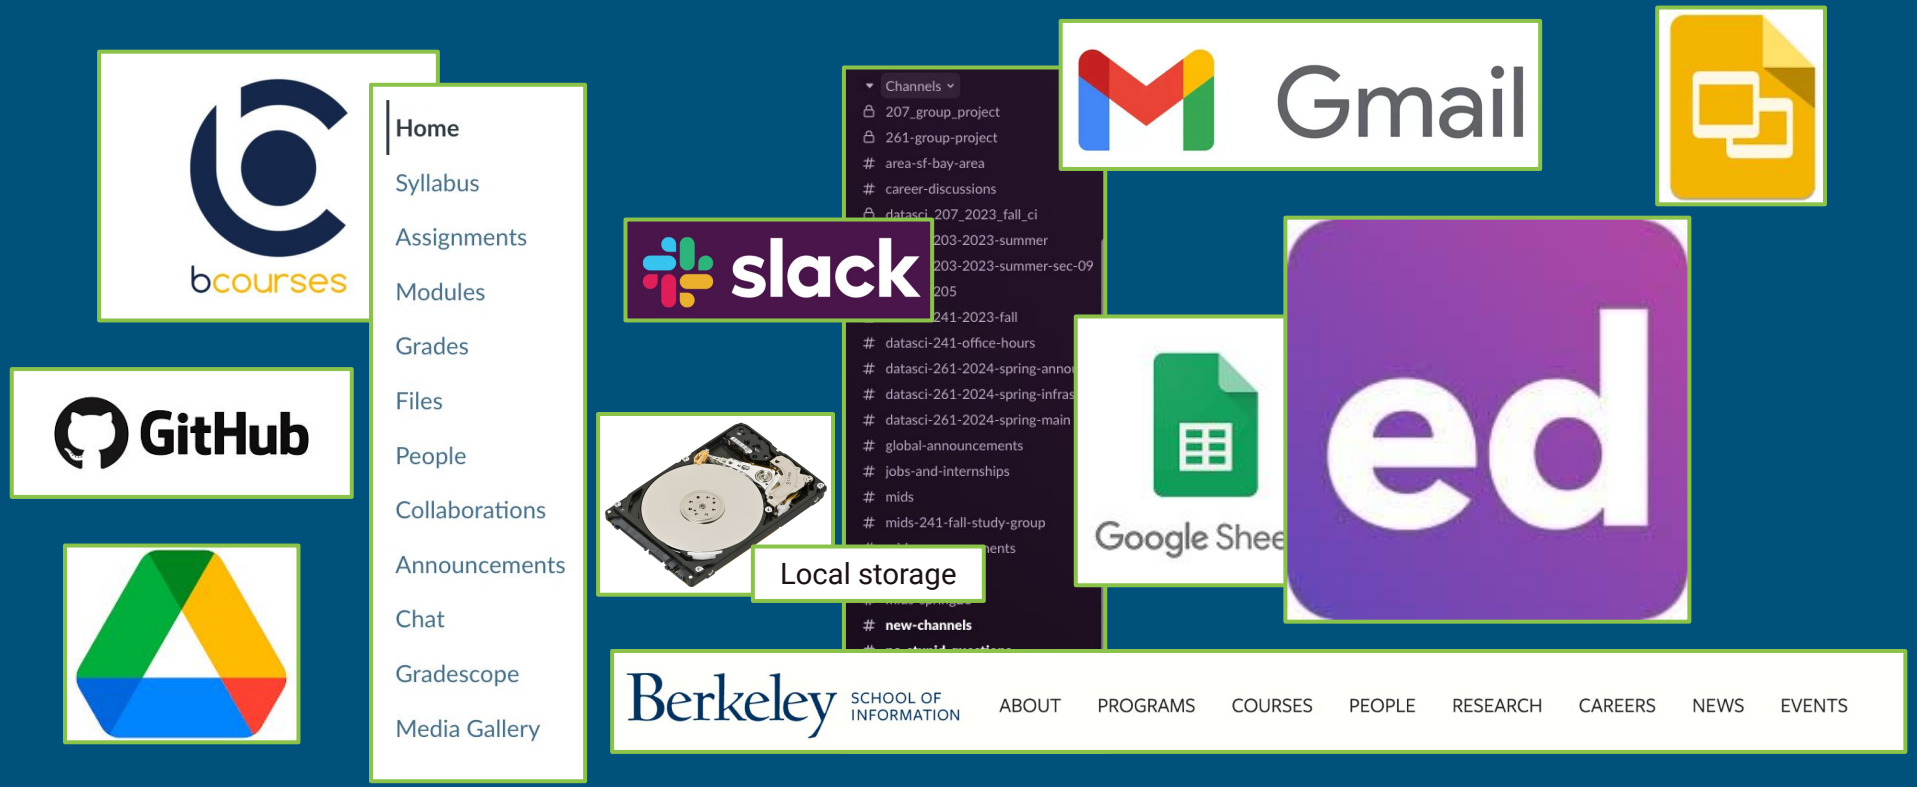

Market research

Knowledge bases / Chat apps
(non-cohesive)

slack

OneDrive

ATLASSIAN

GitHub

G Suite

Discord

Local storage

Cohesive-ish apps

Microsoft Copilot

Recall function
(releasing 6/18/24!)

glue.

Market size

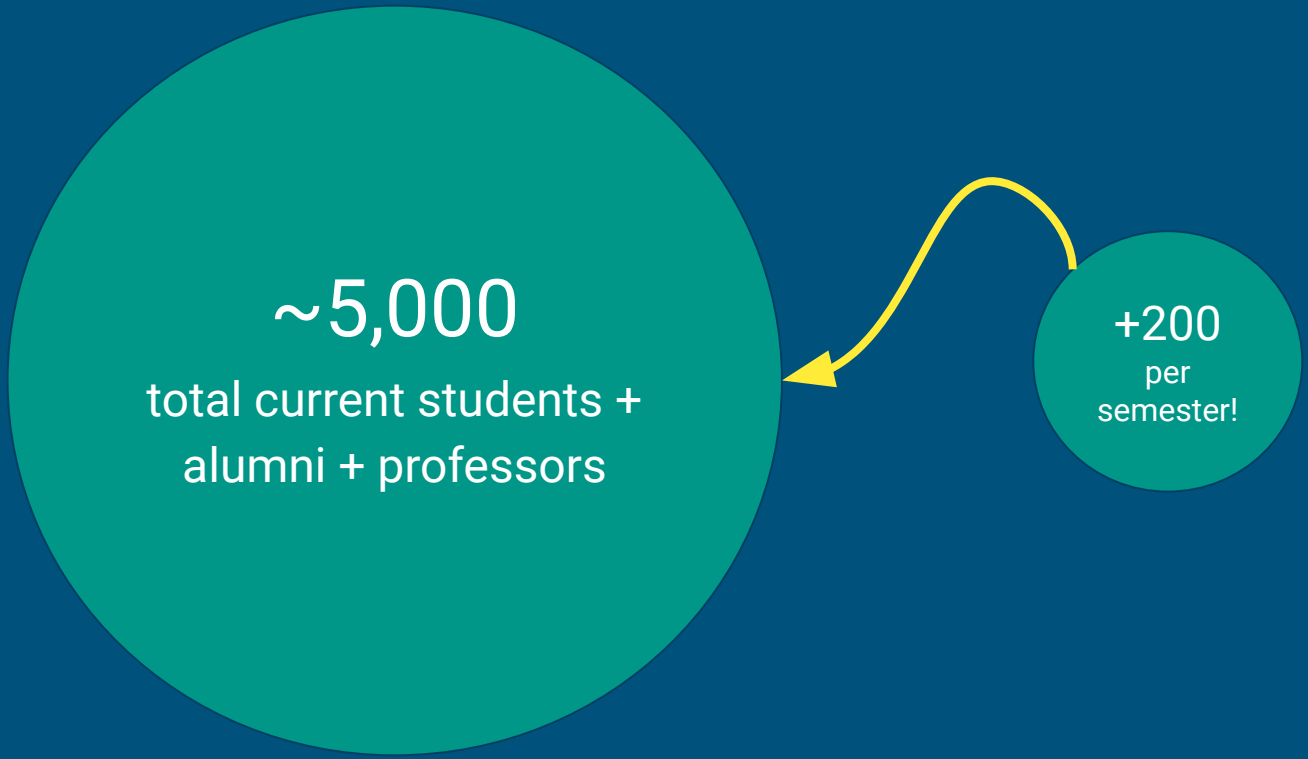
~5,000
total current students +
alumni + professors
+200
per
semester!

What has been tried before

- Previous capstone:
  MIDS Applicant Chatbot (Spring 2021)
- Slack bot 👎

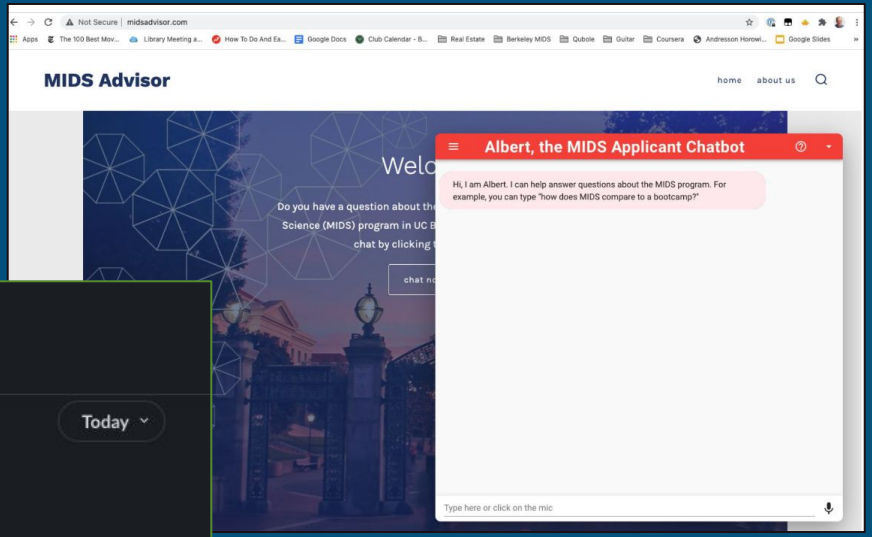

Our Solution and Improvements

RAG based system : Retrieval-Augmented generation is the process of optimizing the output of a large/small language model, so it references an authoritative knowledge base outside of its training data sources before generating a response.

Small Language Model : similar to Large language Model but smaller

User type and customers : Students and Instructors.

Goal is to Reduce implementation and training costs, improve reliability, improve security

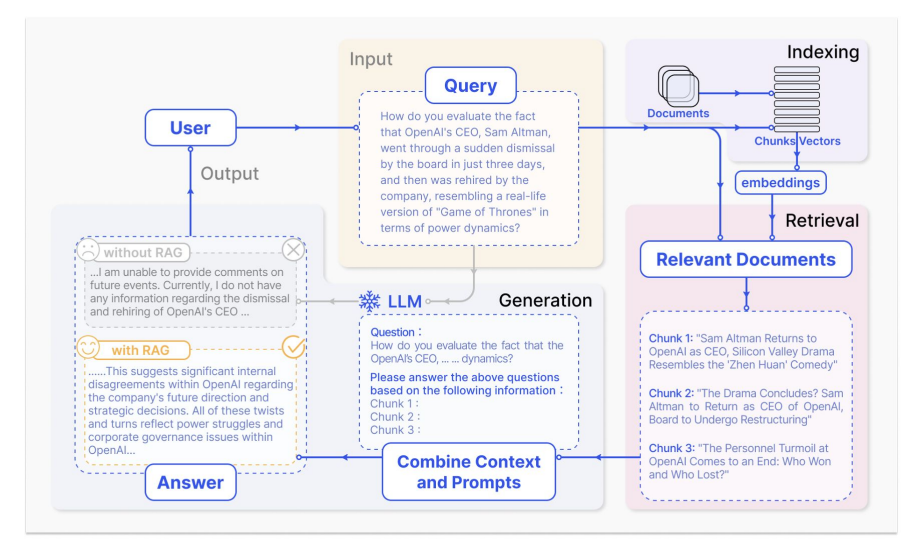

Deep Dive into RAG components

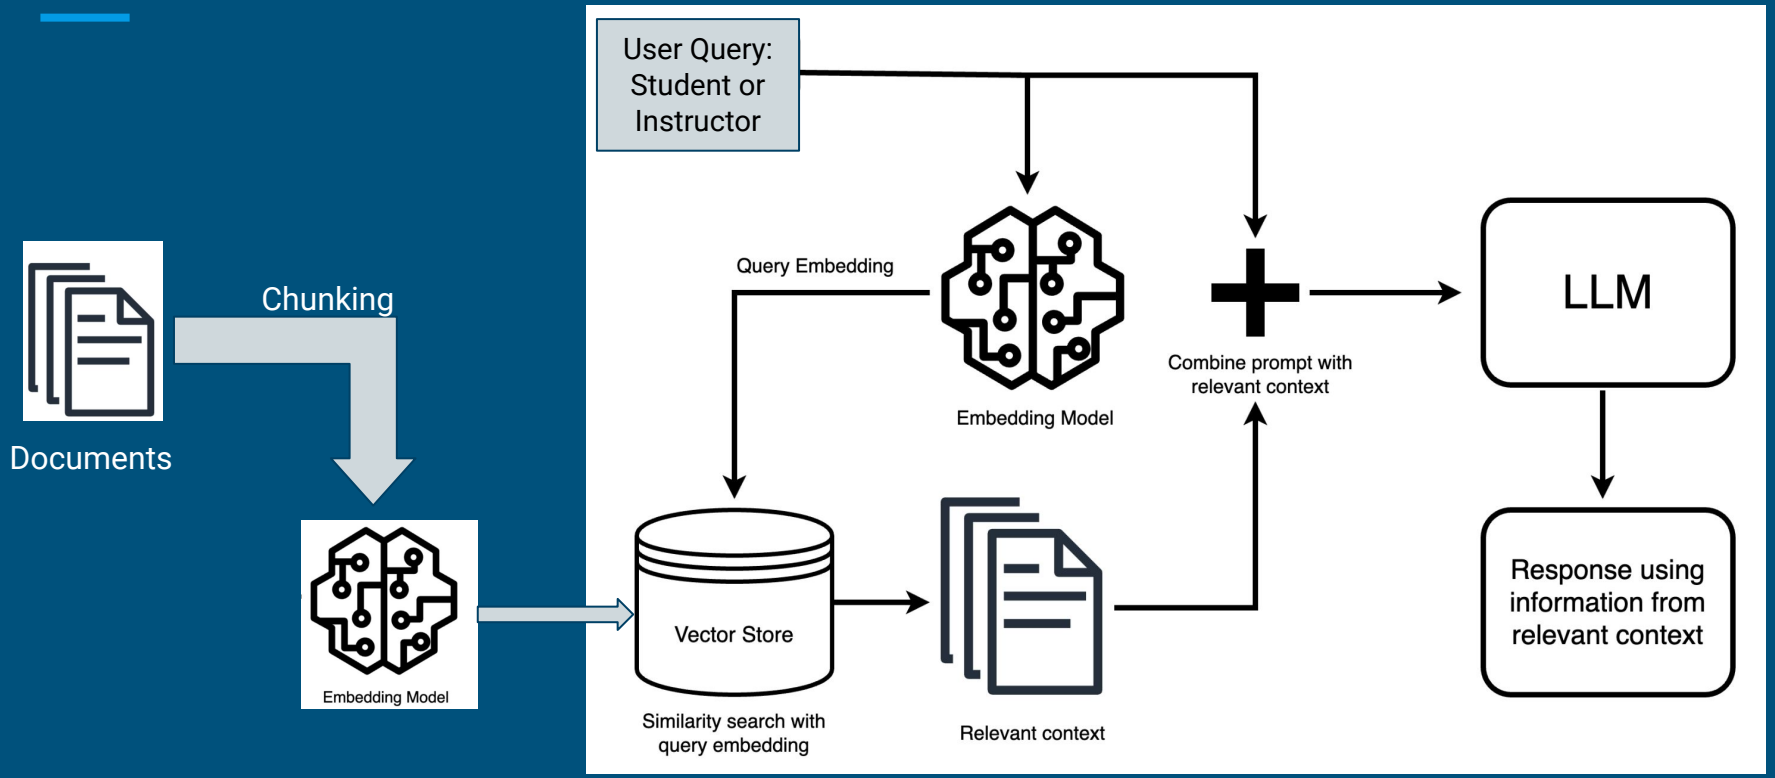

Roadmap - Yes, it's week 5 already

	Week 7	Week 10	Week 12	Week 14
Language Model	Baseline LLM selected SLM Settings Default settings implemented	LoRA User Based Prompting Model Compression Techniques	Prompt Tuning Hyperparameter Tuning	API Chaining
Retrieval	Vector Database Establish Defaults	Search Strategies Reranker	Dynamic Chunking Hybrid Vector Database	Establish Protected Documents
Data	Collection EDA Survey Users	Generate User Specific Responses PDF Processing	Applying additional datasets	
DevOps	AWS Initialization	UI Prototype Security Audit	Gather UI Feedback	Host or Containerize
Evaluation	Research KM KPIs Research NLP Metrics	Apply NLP Experiments	Automated Testing	Human Preference (DPO)

MVP

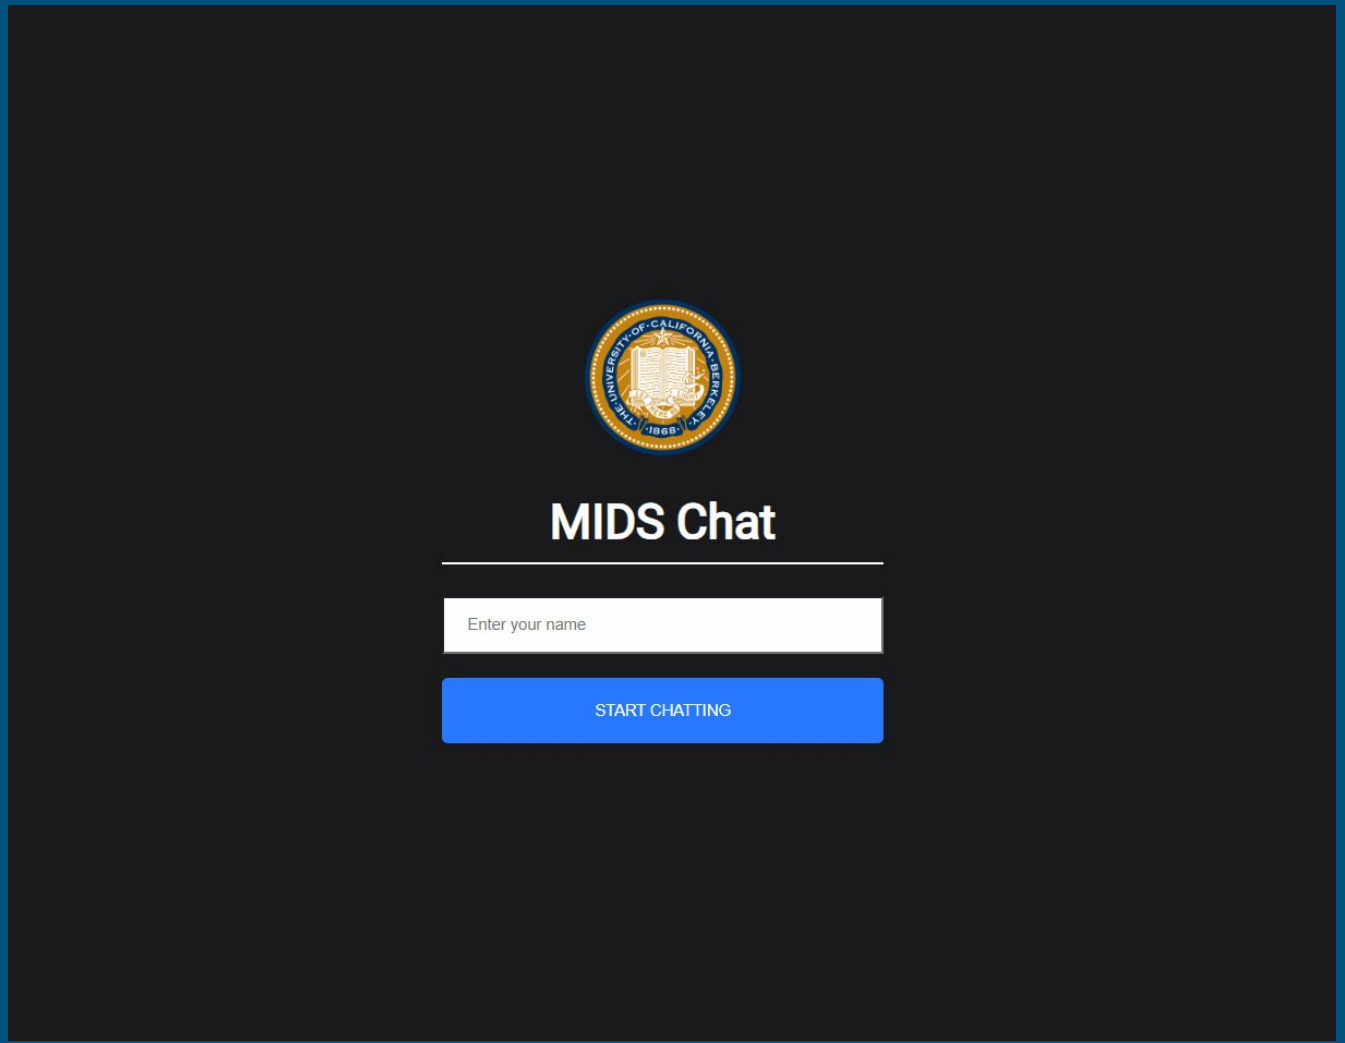

MVP Transition to Real Deal

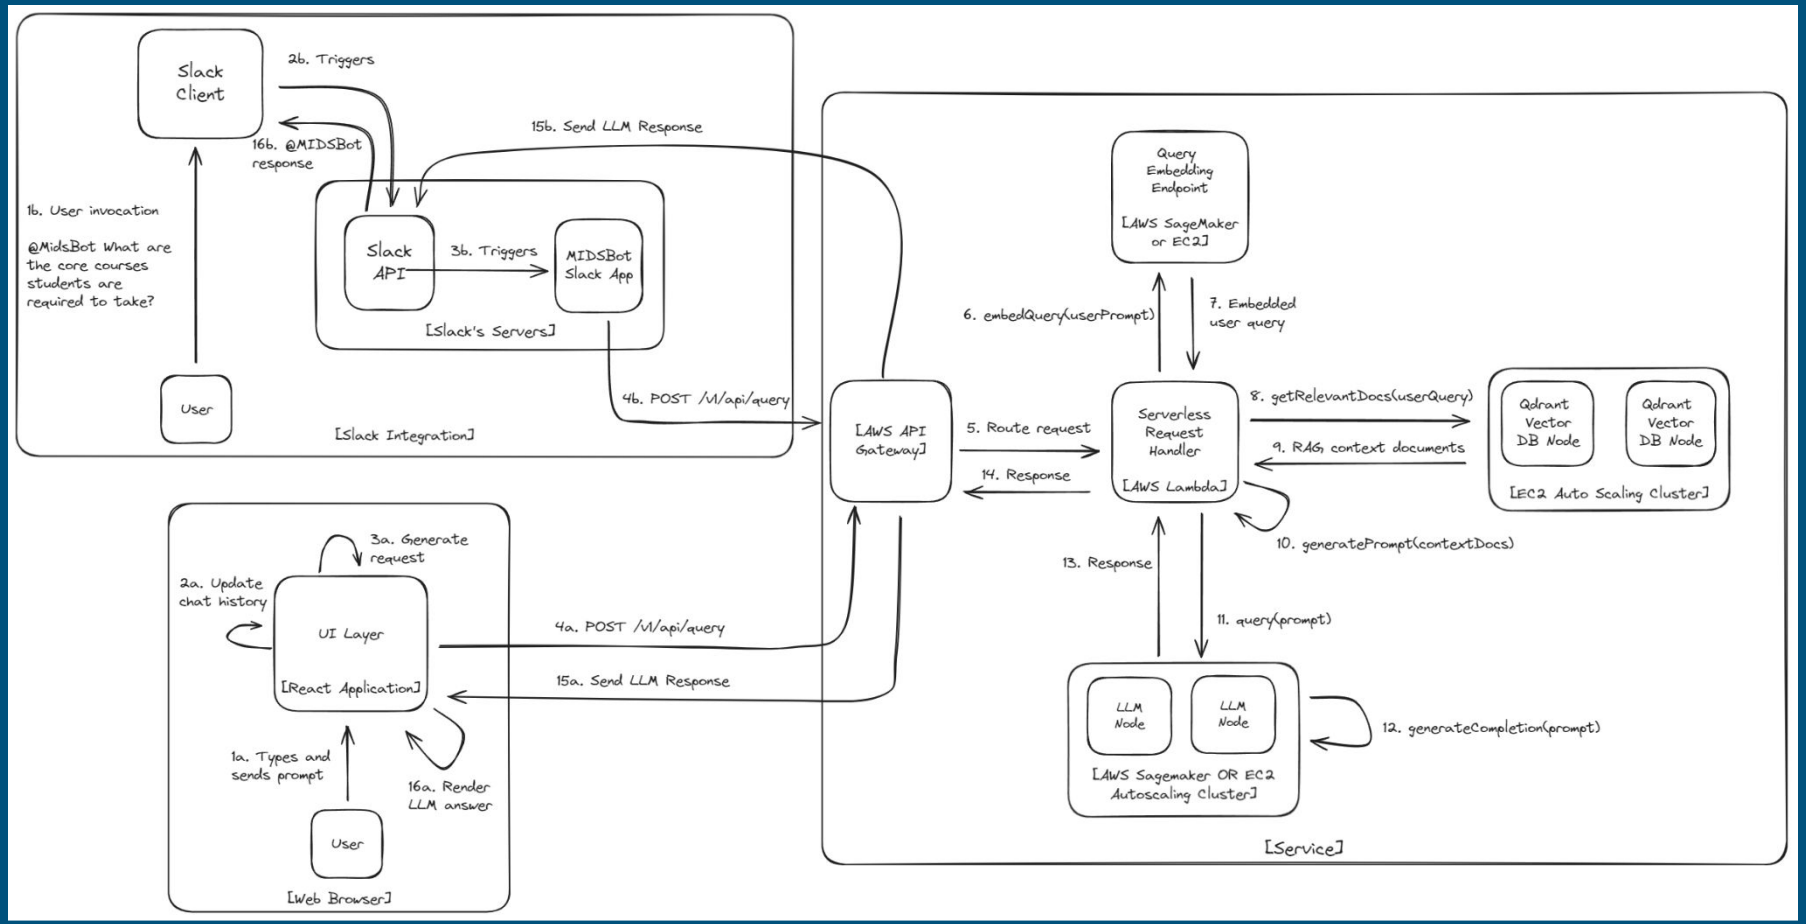

Datasets

Data Collection

- **MIDS Intranet:** BeautifulSoup text scraping
- **Bcourses:** PDF processing MIDS Syllabi
- **Other publicly available sites** as needed
- **Slack**: admin export of #mids-class-rec channel
- **LLM Generated Question Answer Pairs**
  - Google Forms surveying students
  - Using openAi to generate questions AND answers based on available documents

Data Preparation

Data Cleansing

- Remove duplicate content (e.g., repetitive links to Facebook)
- Purge unnecessary content to maintain data relevance and reduce noise

Purge or anonymize all personal information attributes

Data Partitioning by

- Secure pages
- Public pages
- Course-specific pages

User Requirements:

Gather user input to understand their requirements, confirm our assumptions about features, user interface elements,

system usage and system functionality

- Google form
- Open Alpha and Beta to small group of students
- Human feedback via Direct Preference Optimization (DPO)

EDA

N-gram

Named Entity Recognition

Topic Modeling (LDA)

Topic Complexity

Sentiment Analysis

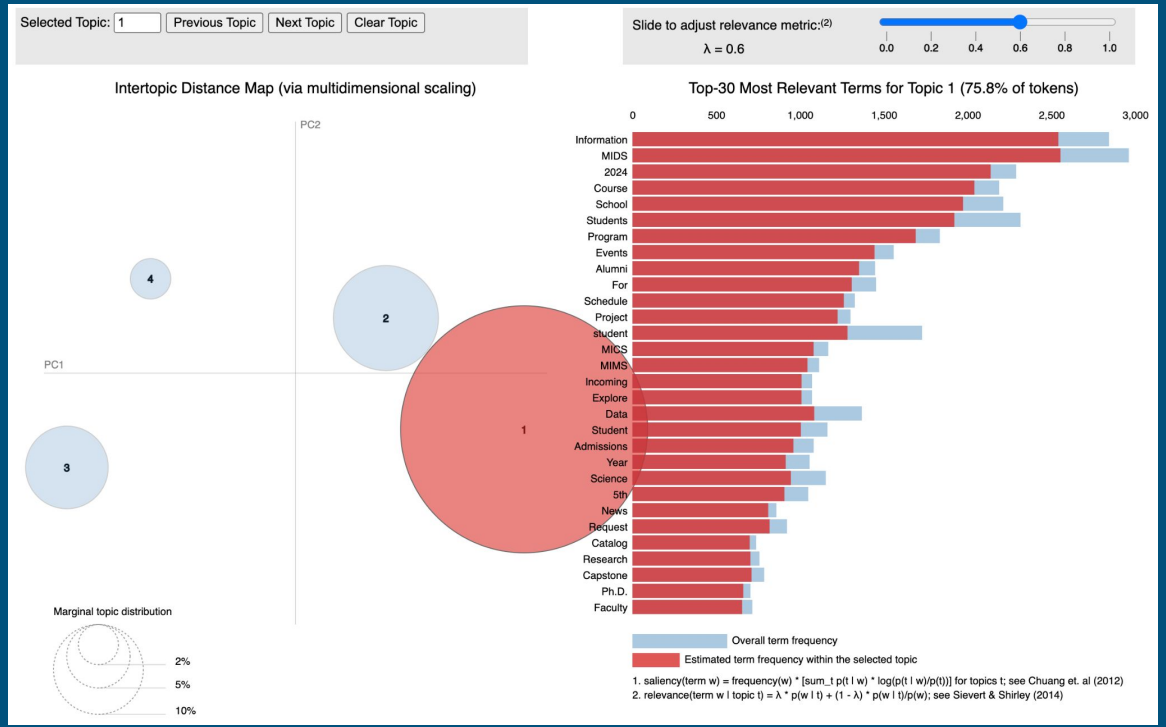

Ethics and Privacy

- Use of Slack
  - Terms and Conditions
- LLM considerations
  - Fine tuning on unbalanced dataset
  - Adversarial attacks
  - Guardrails
- Confidential data
  - PII
  - User sensitive information
  - User Privilege Classification on documents for access

Open Questions and Challenges

- RAG related
  - Generating Reference Answers
  - Model evaluation methods
  - Accessing Slack Data
- Deployment
  - MVP will be hosted locally with a static database
  - Can we build a generative system that run locally for an average users?
- A lot of cool ideas, what are the core features that make the cut?
  - Where will we draw the line ?
  - Not everyone can make varsity

Conclusion

Knowledge is found in many unstructured sources, we unlock it through a RAG-based approach.

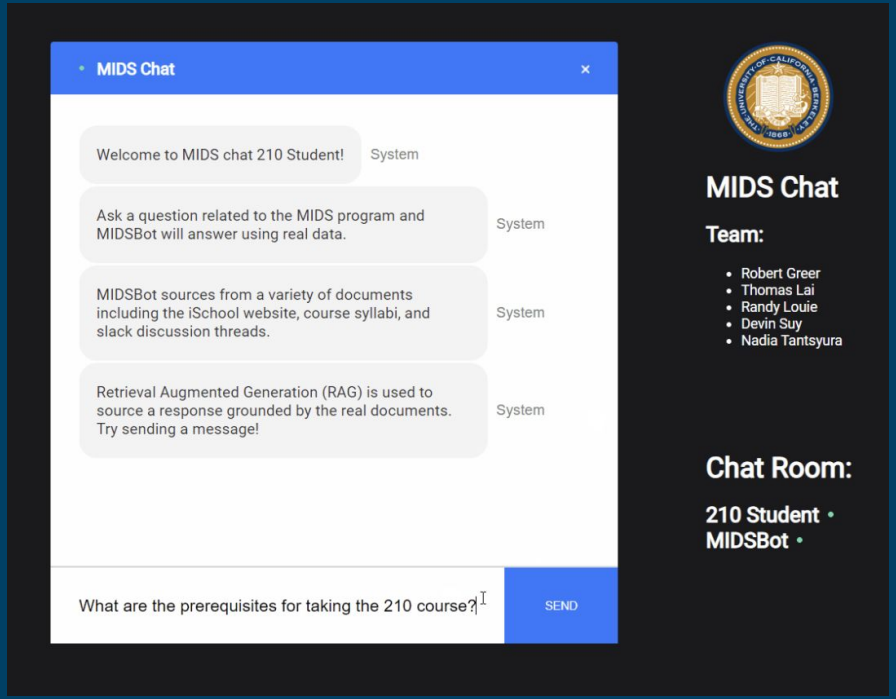

Notes

Pitching - do we keep SLM? Does that align with BErkeley

- Yes, it minimizes infrastructure and implementation complexity for Berkeley.
-